

Introduction to



<https://keras.io/>

What is keras?

“Keras is a high-level neural networks API, written in Python and capable of running on top of TensorFlow, CNTK, or Theano. It was developed with a focus on enabling fast experimentation. Being able to go from idea to result with the least possible delay is key to doing good research. Use Keras if you need a deep learning library that:

- Allows for easy and fast prototyping (through user friendliness, modularity, and extensibility).
- Supports both convolutional networks and recurrent networks, as well as combinations of the two.
- Runs seamlessly on CPU and GPU.”

Keras program structure

```
1 import keras
2 from keras.datasets import cifar10
3 from keras import layers
4 from sklearn import cross_validation
5 from sklearn.utils import shuffle
6 from sklearn.metrics import roc_auc_score
7 from sklearn.preprocessing import OneHotEncoder
8
9 N_EPOCHS= 20
10 BATCH_SIZE= 32
11 LEARNING_RATE= 1e-2 #PLAY with learning rate. try 1e-1, 1e-2 ...
12 N_HIDDEN= 64
13
14 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
15 #Rescale data.
16 x_train= x_train /255.0
17 x_test= x_test /255.0
18 print(x_train.shape, y_train.shape, x_train.max(), x_train.min() )
19 #One-hot-encode labels
20 oh_encoder= OneHotEncoder(n_values=10, sparse=False)
21 y_train= oh_encoder.fit_transform(y_train)
22 y_test= oh_encoder.transform(y_test)
23 print(x_train.shape, y_train.shape)
24
25 print('data ready')
26 #Model definition
27
28 model = keras.models.Sequential()
29 model.add( layers.Conv2D(filters=32, kernel_size= 5, strides=1,padding='same',
30                         activation='relu', input_shape= x_train.shape[1:] ))
31 model.add( layers.MaxPooling2D(pool_size=4, strides=2 ) )
32 model.add( layers.Flatten() )
33 model.add( layers.Dense(N_HIDDEN, activation='relu'))
34 model.add( layers.Dense(y_train.shape[-1], activation='softmax'))
35 sgd= keras.optimizers.SGD(lr=LEARNING_RATE, decay=0., momentum=0., nesterov=False)
36 #Model compilation or loading
37 model.compile(loss='categorical_crossentropy',optimizer=sgd, metrics=[ 'accuracy'])
38 print("network created")
39 print("training begin")
40 #Model training
41 model.fit(x_train, y_train, epochs=N_EPOCHS, batch_size=BATCH_SIZE, validation_split=0.1)
42 #Model evaluation
43 score = tuple(model.evaluate(x_test, y_test, batch_size=BATCH_SIZE))
44 print("Testing evaluation loss %f acc %f roc_auc %f"%score)
45
```

./2_keras_intro/examples/e1_Sequential_keras.py

Data preparation

Model definition

training

evaluation

Data preparation

- If data fits in memory, load it as numpy arrays and train, evaluate or use your model with `model.fit()`, `model.evaluate()`, `model.predict()`.

```
14 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
15 #Rescale data.
16 x_train= x_train /255.0
17 x_test= x_test /255.0
18 print(x_train.shape, y_train.shape, x_train.max(), x_train.min() )
19 #One-hot-encode labels
20 oh_encoder= OneHotEncoder(n_values=10, sparse=False)
21 y_train= oh_encoder.fit_transform(y_train)
22 y_test= oh_encoder.transform(y_test)
23 print(x_train.shape, y_train.shape)
24
25 print("data ready")
40 #Model training
41 model.fit(x_train, y_train, epochs=N_EPOCHS, batch_size=BATCH_SIZE, validation_split=0.1)
42 #Model evaluation
43 score = tuple(model.evaluate(x_test, y_test, batch_size=BATCH_SIZE))
44 print("Testing evaluation loss %f acc %f roc_auc %f"%score)
```

- Otherwise, define a python generator that produces batches or a `keras.utils.Sequence` class and train, evaluate or use your model with `model.fit_generator()`, `model.evaluate_generator()`, `model.predict_generator()`.
- If you want to train on batches as in tensorflow, use `model.train_on_batch()`, `model.test_on_batch()`, `model.predict_on_batch()`

Using generators or keras.utils.Sequence

- Generators, or keras.utils.Sequence are handled by keras very efficiently by means of buffers of batches and multiple CPUs

```
16 (x_train, y_train), (x_test, y_test) = cifar10.load_data()
17 #Rescale data.
18 x_train= x_train /255.0
19 x_test= x_test /255.0
20 print(x_train.shape, y_train.shape, x_train.max(), x_train.min() )
21 #One-hot-encode labels
22 oh_encoder= OneHotEncoder(n_values=10, sparse=False)
23 y_train= oh_encoder.fit_transform(y_train)
24 y_test= oh_encoder.transform(y_test)
25
26 x_train, x_validation, y_train, y_validation = cross_validation.train_test_split(
27                                     x_train, y_train, test_size=0.1, random_state=123)
28
29 DATA_MANAGER_OPTION=1 #Can be 1 or 2
30
31 #OPTION 1
32 def generateBatches(data, labels):
33     ...
34     This is a dummy example, here you will be loading files,
35     applying transformations, etc.
36
37 while True:
38     data, labels= shuffle(data, labels)
39     for i in range(0, data.shape[0], BATCH_SIZE):
40         batch_x= data[i:i+BATCH_SIZE, ...]
41         #transform batch_x or apply data augmentation here
42         batch_y= labels[i:i+BATCH_SIZE, ...]
43         yield batch_x, batch_y
44     #The following line is a possible real example in which data are filenames
45     #return np.array([ resize(imread(file_name), (200, 200))
46     #                  for file_name in batch_x]), np.array(batch_y)
47
48 trainGenerator= generateBatches(x_train, y_train)
49 nBatchesPerEpochTrain= x_train.shape[0]//BATCH_SIZE
50 validationGenerator= generateBatches(x_validation, y_validation)
51 nBatchesPerEpochValid= x_validation.shape[0]//BATCH_SIZE
52 testGenerator= generateBatches(x_test, y_test)
53 nBatchesPerEpochTest= x_test.shape[0]//BATCH_SIZE
54
55 #OPTION 2
56 class myDataManger(Sequence):
57     def __init__(self, x_set, y_set, batch_size= BATCH_SIZE):
58         self.x, self.y = x_set, y_set
59         self.batch_size = batch_size
60
61     def __len__(self):
62         return int(np.ceil(len(self.x) / float(self.batch_size)))
63
64     def __getitem__(self, idx):
65         ...
66         idx is the index of a given example
67
68         batch_x = self.x[idx * self.batch_size:(idx + 1) * self.batch_size]
69         batch_y = self.y[idx * self.batch_size:(idx + 1) * self.batch_size]
70         return batch_x, batch_y
71
72         ...
73
74         ...
75
76         ...
77
78         ...
79
79 #Model training
80 if DATA_MANAGER_OPTION==1:
81     #OPTION 1
82     model.fit_generator(trainGenerator, steps_per_epoch=nBatchesPerEpochTrain, epochs=N_EPOCHS,
83                         validation_data= validationGenerator, validation_steps=nBatchesPerEpochValid,
84                         workers=2, max_queue_size=10, use_multiprocessing=True )
85 if DATA_MANAGER_OPTION==2:
86     #OPTION 2
87     model.fit_generator(myDataManger(x_train, y_train), epochs=N_EPOCHS,
88                         validation_data= myDataManger(x_validation, y_validation),
89                         workers=2, max_queue_size=10 )
90
91
92
93
94
95
96
97
98
99
100
```

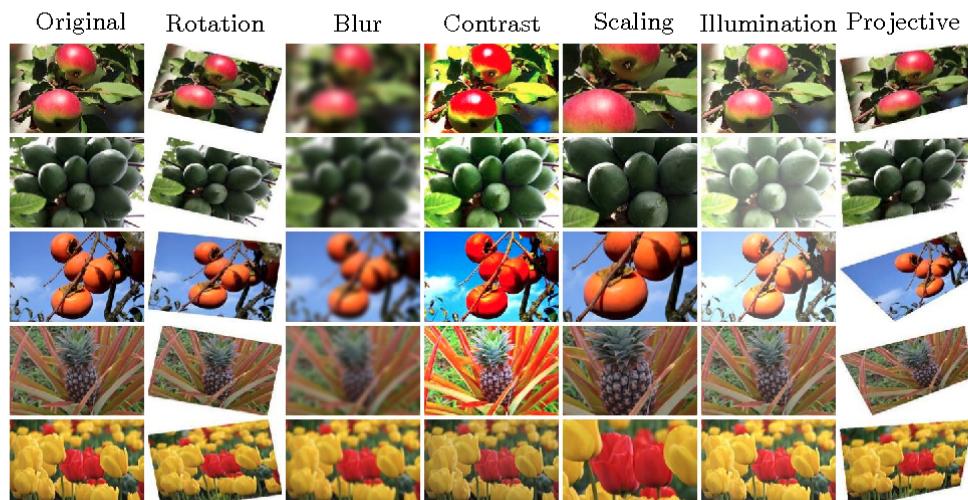
Data augmentation

./2_keras_intro/examples/e2_b_dataAugmentation.py

For images or sequences you can use built-in classes.

<https://keras.io/preprocessing/image/#imagedatagenerator-methods>

```
30
31 if AUGMENTATION_OPTION==1:
32     #OPTION 1: AUTOMATIC AUGMENTATION. Data need to fit in memory or being stored in one directory per class
33     #Image augmentation
34     dataAugmentator = ImageDataGenerator(rotation_range=10.0, horizontal_flip=True, fill_mode='reflect')
-- 
112
113 if AUGMENTATION_OPTION==1:
114     #When data is located as png, jpg, bmp, files ImageDataGenerator can use then directly with
115     #dataAugmentator.flow_from_directory
116     model.fit_generator(dataAugmentator.flow(x_train, y_train), epochs=N_EPOCHS,
117                         validation_data= (x_validation, y_validation),
118                         workers=2, max_queue_size=10, use_multiprocessing=True )
```



Data augmentation

For other types of data you have to handle it on your way.

- Easy when coupled with generators or keras.utils.Sequence.

```
37 #OPTION 2: MANUAL AUGMENTATION. If you want to use generators
38 from skimage.transform import rotate
39 class myDataManger(Sequence):
40     def __init__(self, x_set, y_set, batch_size= BATCH_SIZE, augment=True):
41         self.x, self.y = x_set, y_set
42         self.batch_size = batch_size
43         if augment:
44             self.augmentBatch= self._augmentBatch
45         else:
46             self.augmentBatch= lambda arg: arg
47
48     def __len__(self):
49         return int(np.ceil(len(self.x) / float(self.batch_size)))
50
51     def _random_flip_leftright(self, batch):
52         for i in range(len(batch)):
53             if bool(random.getrandbits(1)):
54                 batch[i] = np.fliplr(batch[i])
55         return batch
56
57     def _random_flip_updown(self, batch):
58         for i in range(len(batch)):
59             if bool(random.getrandbits(1)):
60                 batch[i] = np.flipud(batch[i])
61         return batch
62
63     def _random_rotation(self, batch, max_angle=15.0):
64         for i in range(len(batch)):
65             if bool(random.getrandbits(1)):
66                 # Random angle
67                 angle = random.uniform(-max_angle, max_angle)
68                 batch[i] = rotate(batch[i], angle, mode="reflect")
69         return batch
70
71     def _augmentBatch(self, batch):
72         if bool(random.getrandbits(1)):
73             batch= self._random_flip_leftright(batch)
74         if bool(random.getrandbits(1)):
75             batch= self._random_rotation(batch, 10.0)
76         return batch
77
78     def __getitem__(self, idx):
79         ...
80             idx is the index of a given example
81         ...
82         batch_x = self.x[idx * self.batch_size:(idx + 1) * self.batch_size]
83         batch_y = self.y[idx * self.batch_size:(idx + 1) * self.batch_size]
84         return self.augmentBatch(batch_x), batch_y
85
```

./2_keras_intro/examples
/e2_b_dataAugmentation
.py

Model definition

Two alternatives:

- Easy models (Feed-forward, no bifurcations, etc.)
 - Use Sequential API
(<https://keras.io/models/sequential/>)
- Complicate models
 - Use Functional API (<https://keras.io/models/model/>)

Sequential API

- keras.models.Sequential class

```
42 model = keras.models.Sequential()
43 model.add( layers.Conv2D(filters=32, kernel_size= 5, strides=1,padding='same',
44                         activation='relu', input_shape= x_train.shape[1:] ))
45 model.add( layers.MaxPooling2D(pool_size=4, strides=2 ) )
46 model.add( layers.Flatten() )
47 model.add( layers.Dense(N_HIDDEN, activation='relu'))
48 model.add( layers.Dense(y_train.shape[-1], activation='softmax'))
49 sgd= keras.optimizers.SGD(lr=LEARNING_RATE, decay=0., momentum=0., nesterov=False)
50
51 model.compile(loss='categorical_crossentropy',optimizer=sgd, metrics=['accuracy', 'roc_auc'])
52 print("network created")
```

<https://keras.io/>

./2_keras_intro/examples/e1_Sequential_keras.py

Functional API

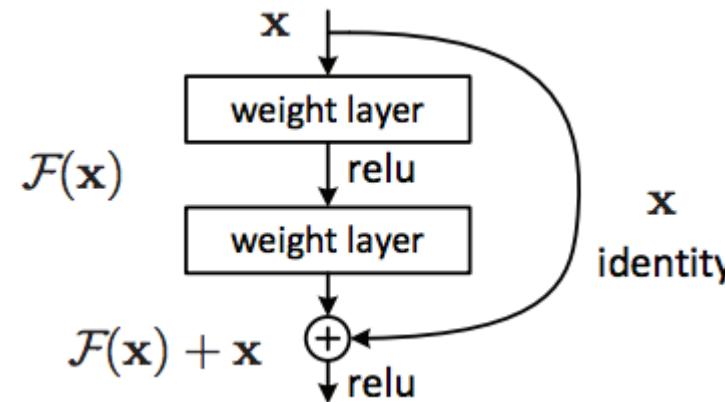
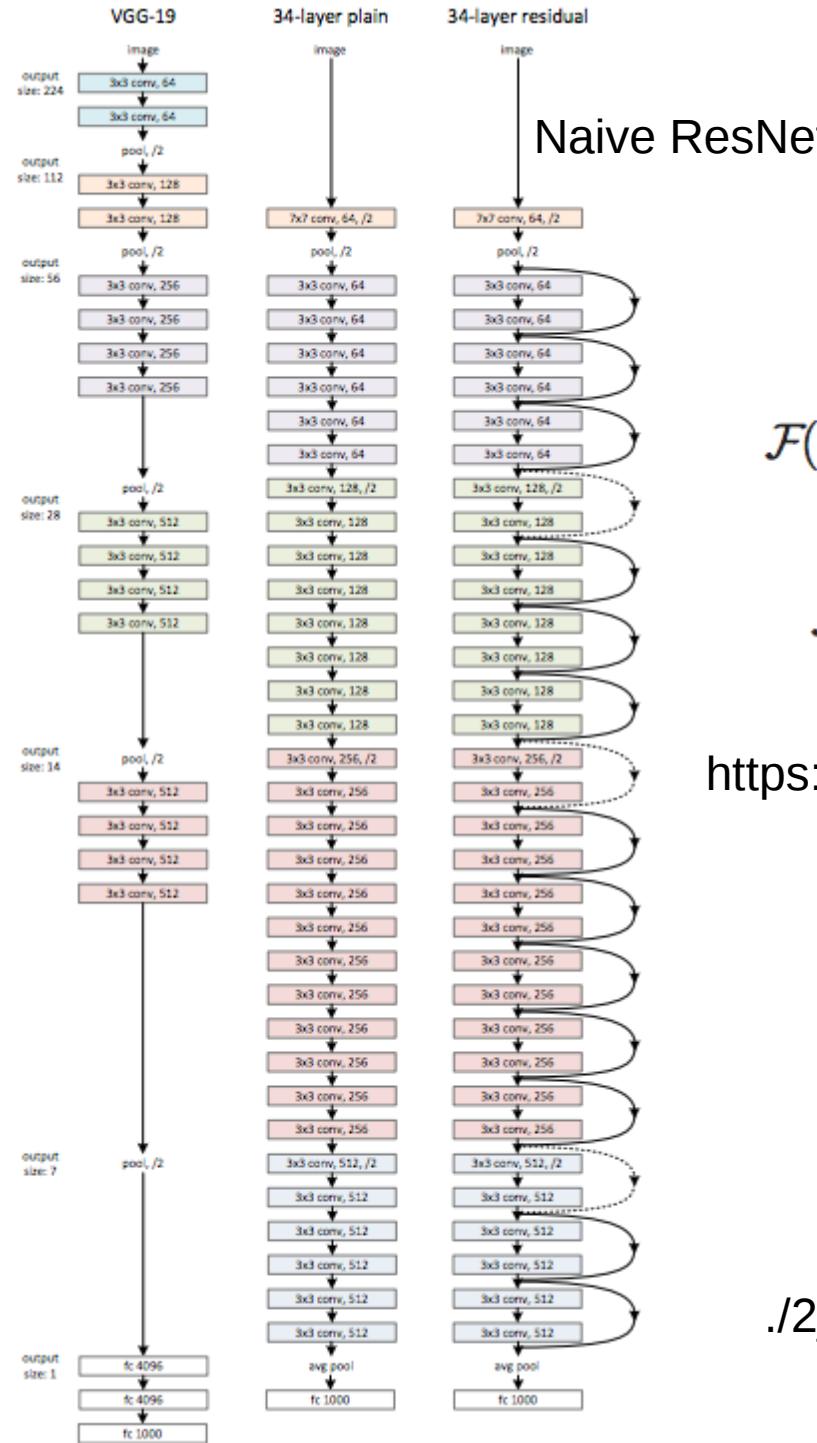
- keras.models.Model class

```
-->
40 #Model definition
41
42 net_input= layers.Input( shape= x_train.shape[1:])
43 net_out= layers.Conv2D(filters=32, kernel_size= 5, strides=1,padding='same',
44                         activation='relu'  )(net_input)
45 net_out= layers.MaxPooling2D(pool_size=4, strides=2 )(net_out)
46 net_out= layers.Flatten()(net_out)
47 net_out= layers.Dense(N_HIDDEN, activation='relu')(net_out)
48 net_out= layers.Dense(y_train.shape[-1], activation='softmax')(net_out)
49
50 model = keras.models.Model( inputs=net_input, outputs= net_out )
51
52 sgd= keras.optimizers.SGD(lr=LEARNING_RATE, decay=0., momentum=0., nesterov=False)
53
54 model.compile(loss='categorical_crossentropy',optimizer=sgd, metrics=[ 'accuracy' , roc_auc])
55 print("network created")
```

<https://keras.io/>

./2_keras_intro/examples/e3_Funltional_keras.py

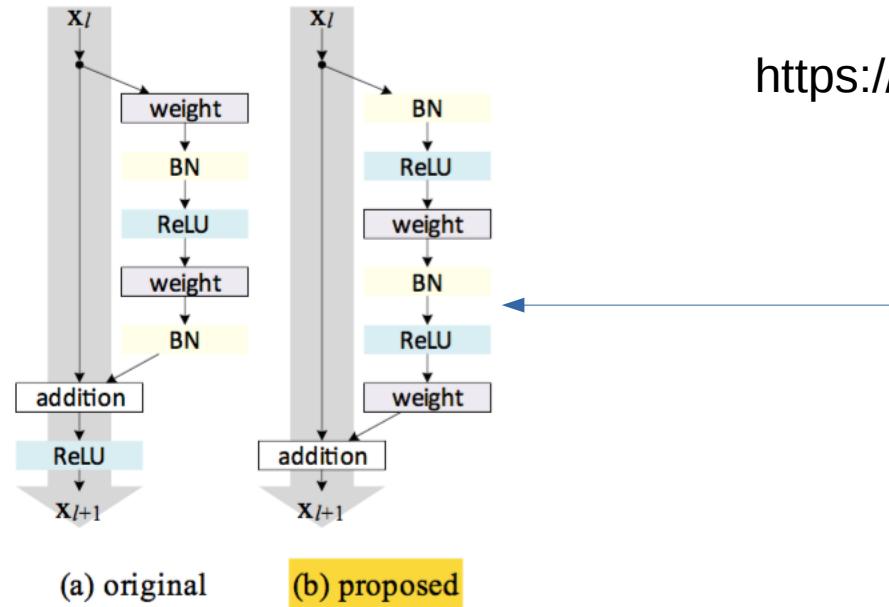
Bifurcation example: Residual Network



<https://arxiv.org/pdf/1512.03385.pdf>

`./2_keras_intro/examples/e4_a_Residual_keras.py`

Bifurcation example: Residual Network



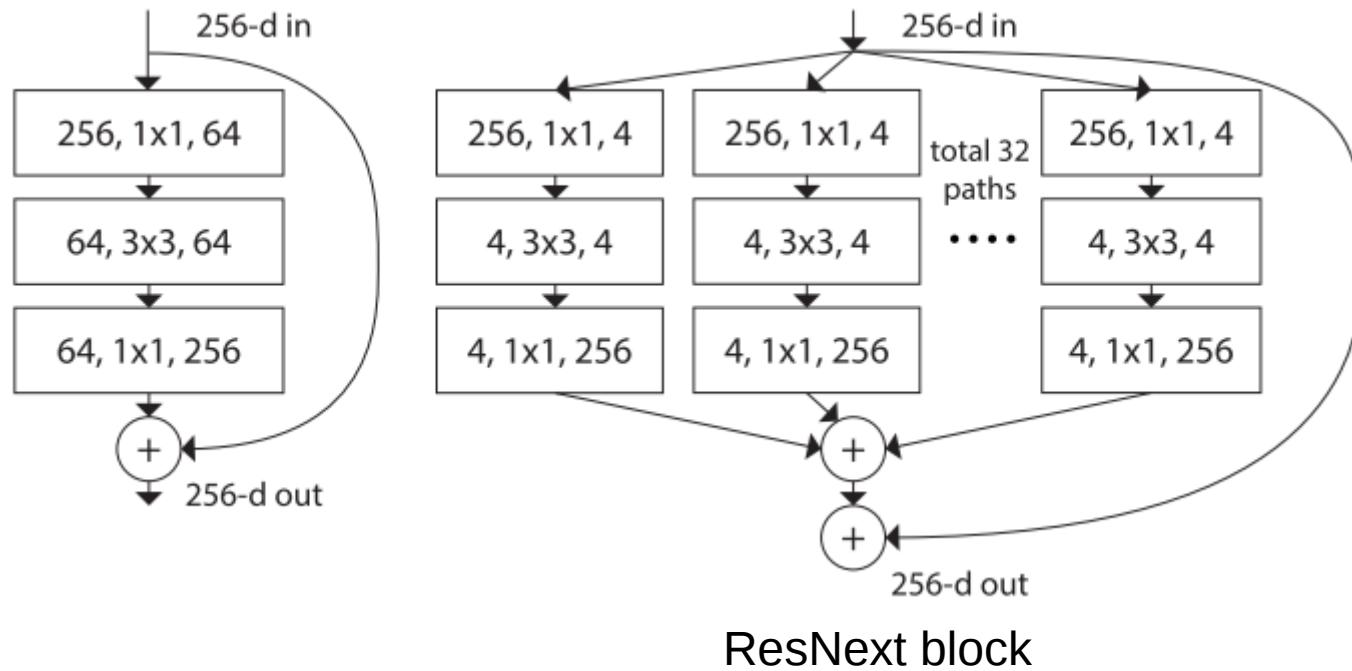
<https://arxiv.org/pdf/1512.03385.pdf>

We will be suffering overfitting:

- Dropout
- Regularization (l1 or l2)
- Data augmentation

Batch Normalization (BN) makes training faster, and it may slightly reduce overfitting.

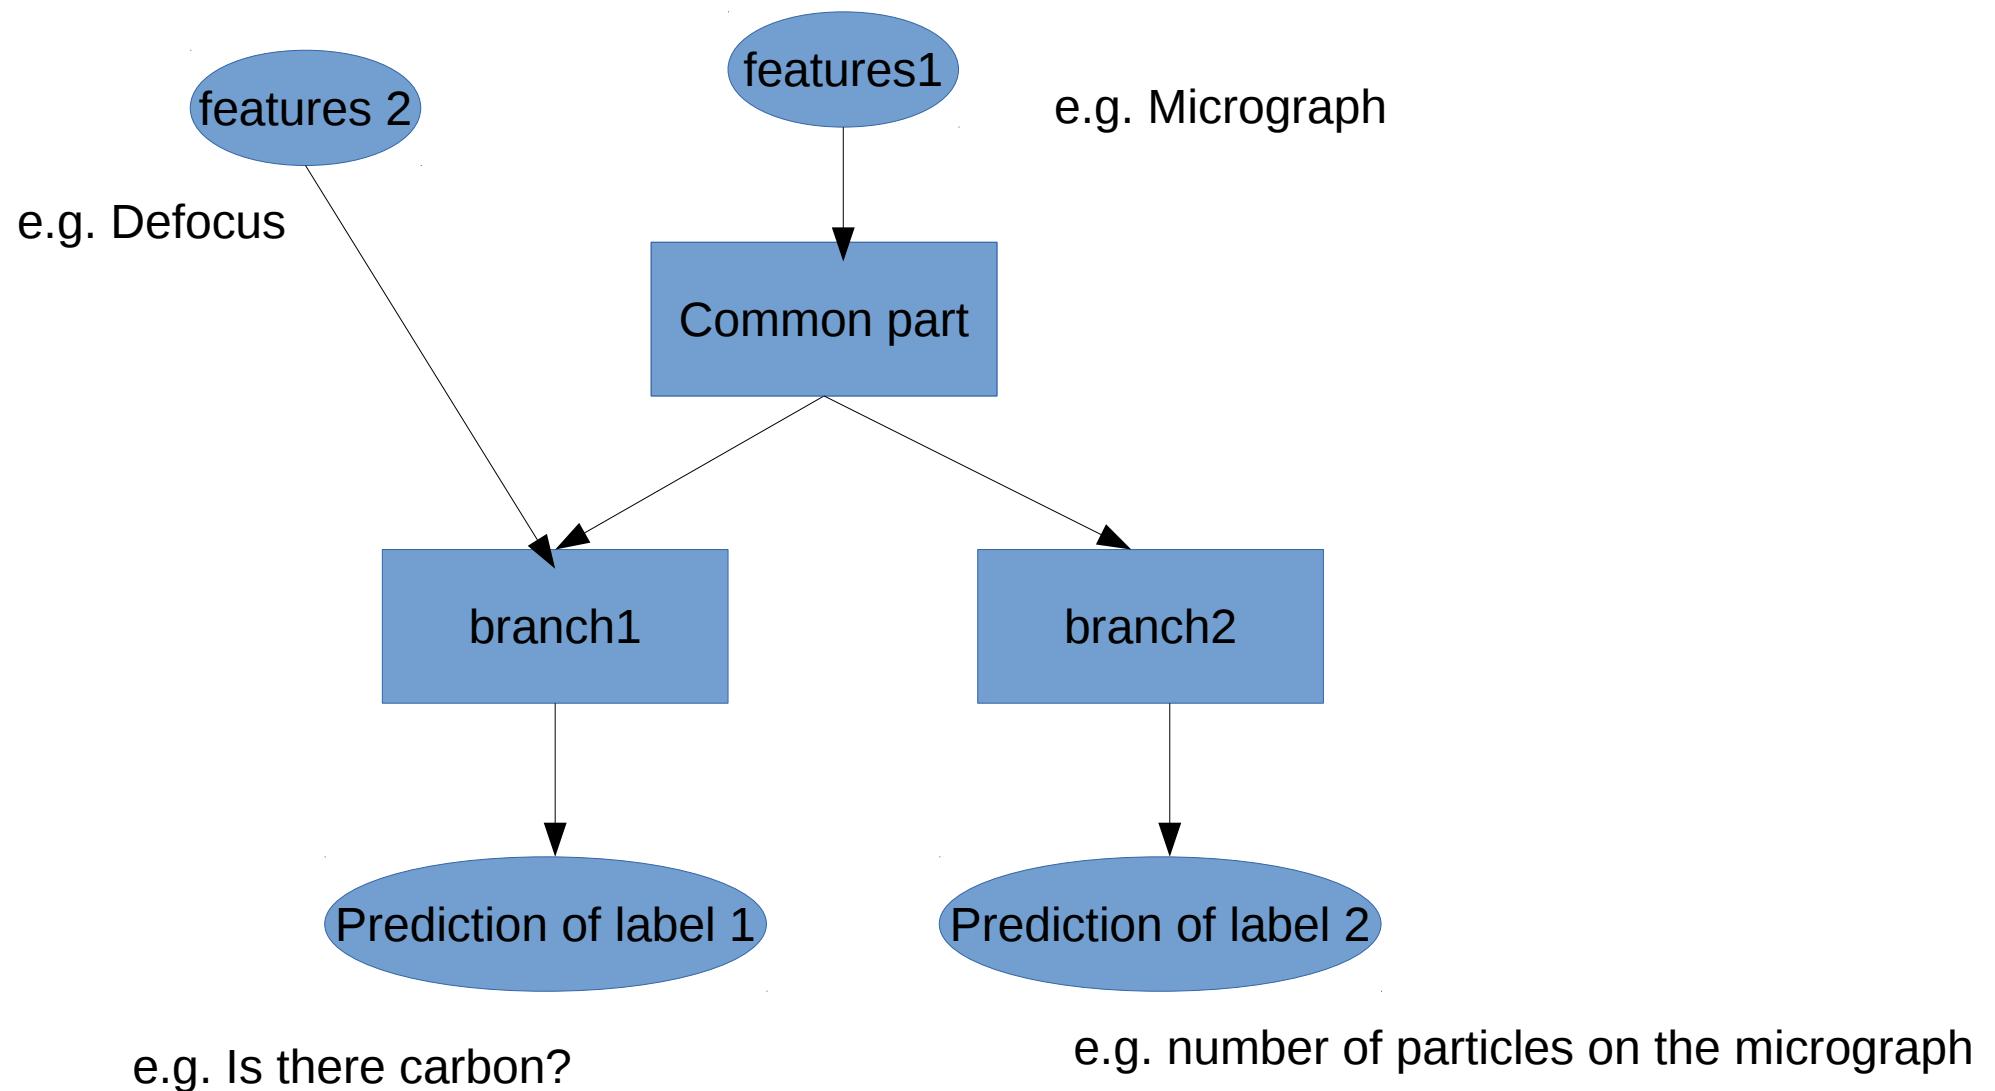
Your turn: ResNext like block



<https://arxiv.org/pdf/1611.05431.pdf>

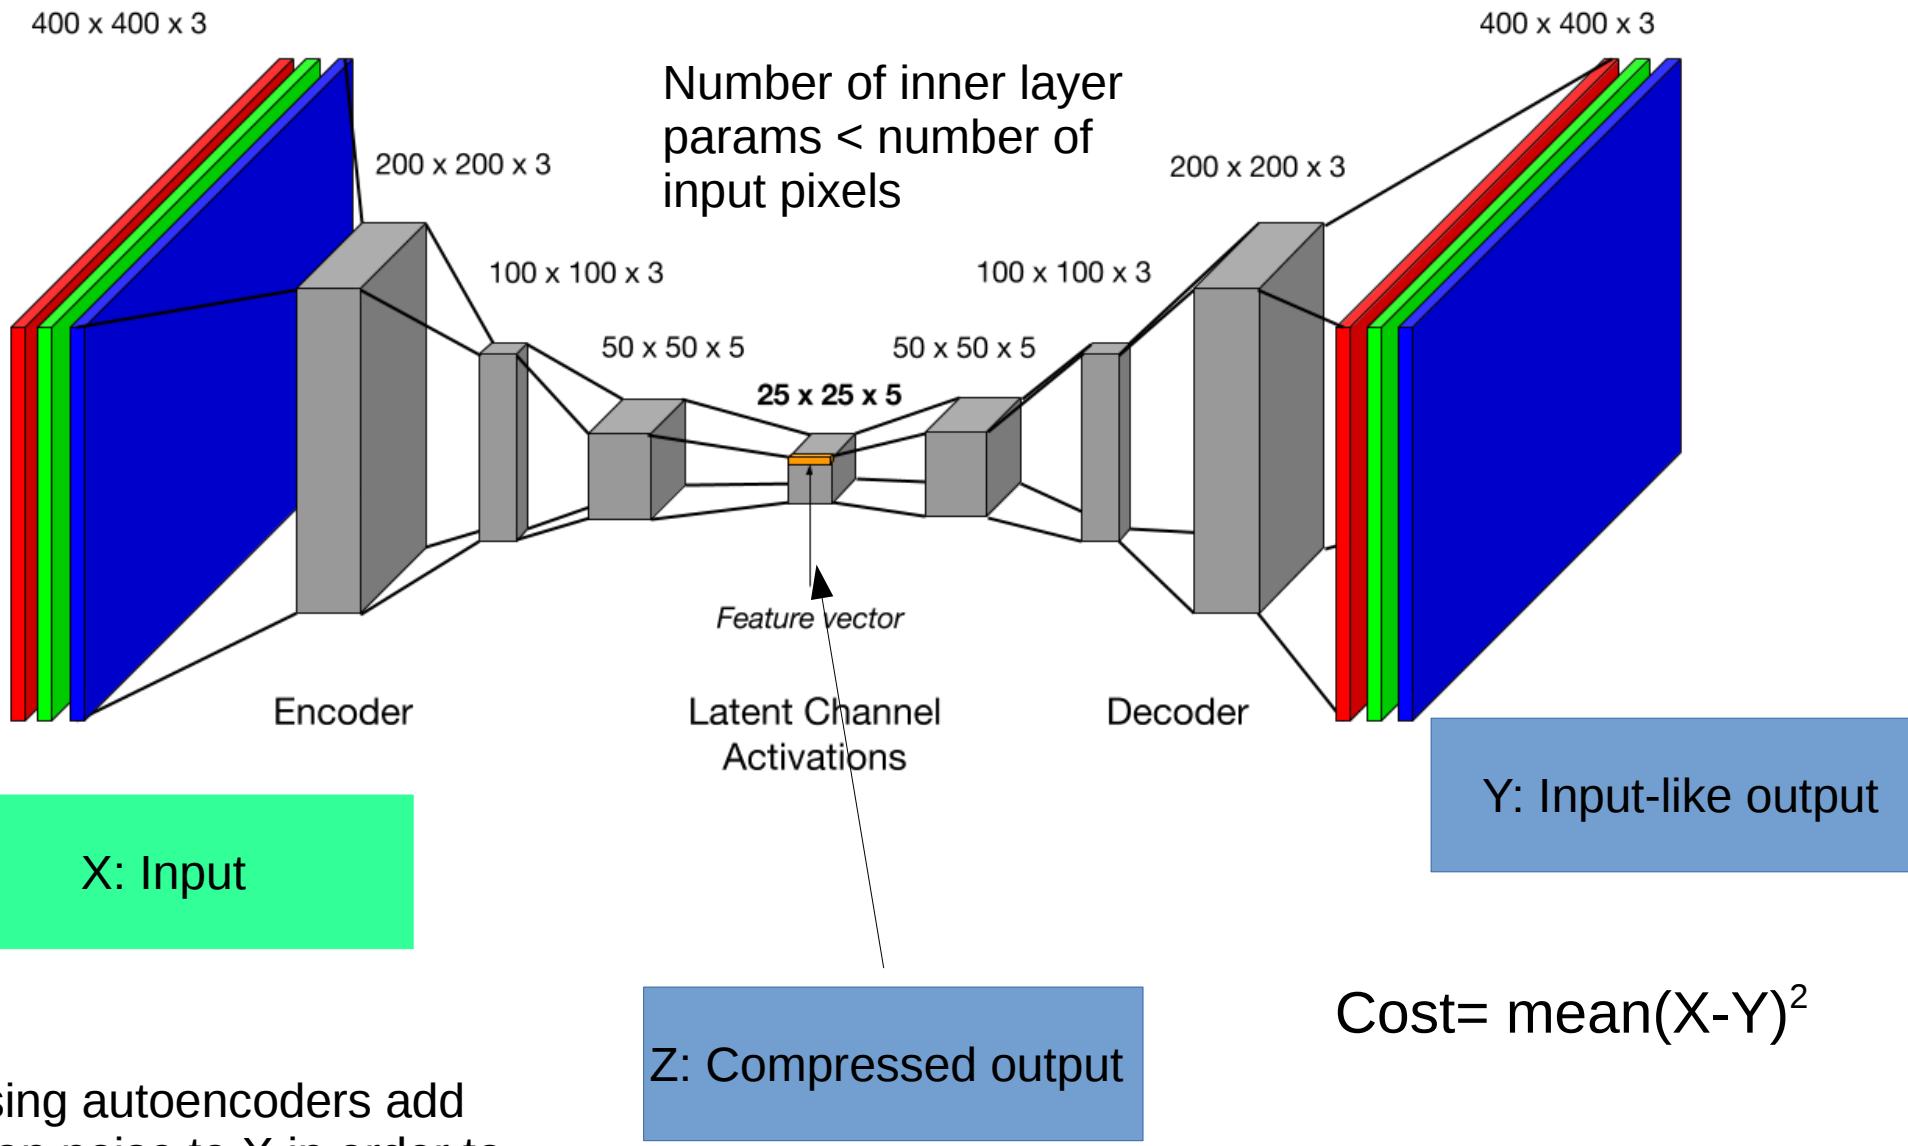
`./2_keras_intro/exercises_sol/sol1_ResNext.py`

Several inputs and/or outputs models



`./2_keras_intro/examples/e5_severalInputsOutputs.py`

Your turn: implement autoencoder



Denoising autoencoders add gaussian noise to X in order to force better representation

`./2_keras_intro/exercises_sol/sol2_convAutoencoder.py`

Using low level keras: keras.backend

- Keras.backend is similar to tf low level calls.
- Used inside Lambda layers and loss or metrics definitions.
 - Direct tf calls can be done inside

```
34 def roc_auc(y_true, y_pred):
35     auc = tf.metrics.auc(y_true, y_pred)[1]
36     keras.backend.get_session().run(tf.local_variables_initializer())
37     return auc
```

```
# add a x -> x^2 layer
model.add(Lambda(lambda x: x ** 2))
```

```
# add a layer that returns the concatenation
# of the positive part of the input and
# the opposite of the negative part

def antirectifier(x):
    x -= K.mean(x, axis=1, keepdims=True)
    x = K.l2_normalize(x, axis=1)
    pos = K.relu(x)
    neg = K.relu(-x)
    return K.concatenate([pos, neg], axis=1)

def antirectifier_output_shape(input_shape):
    shape = list(input_shape)
    assert len(shape) == 2 # only valid for 2D tensors
    shape[-1] *= 2
    return tuple(shape)

model.add(Lambda(antirectifier,
                 output_shape=antirectifier_output_shape))
```

Training tricks: callbacks

Callbacks are functions that are applied at certain stages of training, typically at the end or the beginning of a new epoch. Passed as argument to model.fit()

Very useful default callbacks:

- keras.callbacks.ModelCheckpoint.
- keras.callbacks.EarlyStopping.
- keras.callbacks.ReduceLROnPlateau
- keras.callbacks.TensorBoard

`./2_keras_intro/examples/e6_model_with_callbacks.py`

keras.callbacks.ModelCheckpoint saves keras model after epochs that have improved evaluation metric.

Save model manually

./2_keras_intro/examples/e6_model_with_callbacks.py

```
52
53 from keras.models import load_model
54 model.save('my_model.h5') # creates a HDF5 file 'my_model.h5'
55 del model # deletes the existing model
56 # returns a compiled model
57 # identical to the previous one
58 model = load_model('my_model.h5')
59
60 Checkpointer saving model
61
62
63 from keras.callbacks import ModelCheckpoint
64
65 checkpointer = ModelCheckpoint(filepath='my_model.h5', save_best_only=True)
66 model.fit(x_train, y_train, batch_size=128, epochs=20, verbose=0,
67             validation_data=(X_val, Y_val), callbacks=[checkpointer])
68
69 del model # deletes the existing model
70 model = load_model('my_model.h5')
71
72
73 keras.callbacks.ModelCheckpoint(filepath, monitor='val_loss', verbose=0,
74     save_best_only=False, save_weights_only=False, mode='auto', period=1)
```

ReduceLROnPlateau helps optimizer to improve results by lowering learning rate when stuck

```
keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.1,  
patience=10, verbose=0, mode='auto', min_delta=0.0001, cooldown=0,  
min_lr=0)
```

EarlyStopping prevents overfitting by stopping when error is not improving

```
keras.callbacks.EarlyStopping(monitor='val_loss', min_delta=0, patience=0,  
verbose=0, mode='auto', baseline=None)
```

Tensorboard usage

```
keras.callbacks.TensorBoard(log_dir='./logs', histogram_freq=0, batch_size=32,  
write_graph=True, write_grads=False, write_images=False, embeddings_freq=0,  
embeddings_layer_names=None, embeddings_metadata=None, embeddings_data=None)
```

```
101  
102 tf.summary.image("input_images", model.get_layer("initial_conv").input)  
103 tf.summary.image("first layer output filter0", model.get_layer("initial_conv").output[...,0:1]) #input must be a 4-D tensor with last dimension 1 or 3  
104  
105 callbacksList+= [keras.callbacks.TensorBoard(log_dir='./logs', batch_size=BATCH_SIZE, histogram_freq=1, write_graph=True, write_grads=True, write_images=False)]
```

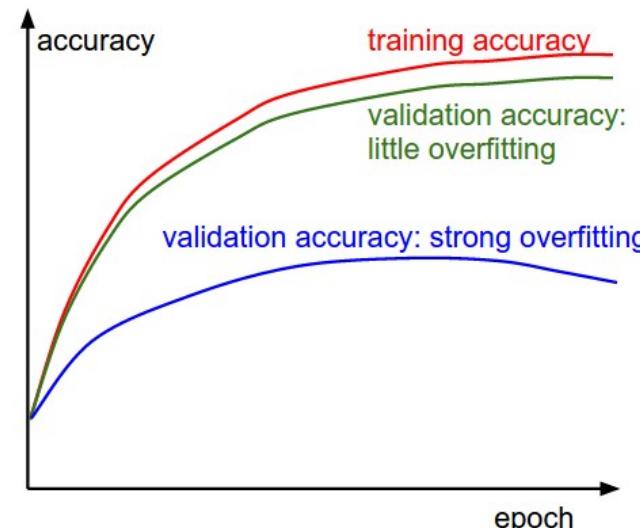
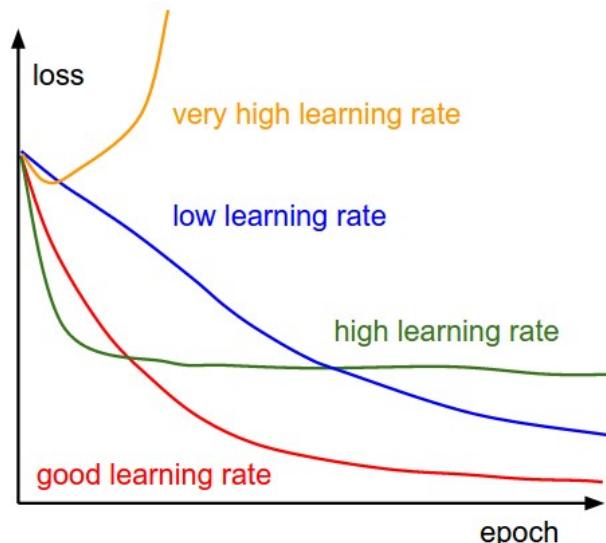
Warning: Makes training slower specially if write_xx=True. Use only for debugging.

If True all layers will be recorded, better with tf.summary

Your turn: do your best on cifar10

My advices:

- Always normalize data
- First try to overfit small subset of training data using big network. If you are not able to do so, there is some problem with your data or your network.
- Use Adam optimizer. If loss oscillates, try smaller learning rate. If keeps vibrating, try different values of epsilon parameter. Increasing batch size also may help.
- Use ReduceLROnPlateau.
- Use checkpoints and employ the one with the best performance. (Sometimes using several checkpoints and averaging them works better).
- Start with a well established architecture.
- Monitor train and validation loss or accuracy. Diagnose whether you are suffering overfitting or underfitting and apply corrective actions (more/less regularization, data augmentation...).



The end